

# “FuN51llyM4ch1n3”

## Official Writeup

27<sup>th</sup> January 2026

**Prepared By:** {CYNX} Team

**Title:** FuN51llyM4ch1n3

**Description:** Beep boop! The age of robots is here! But this one's kinda funky. Maybe it stores some kind of password? Try running it :D

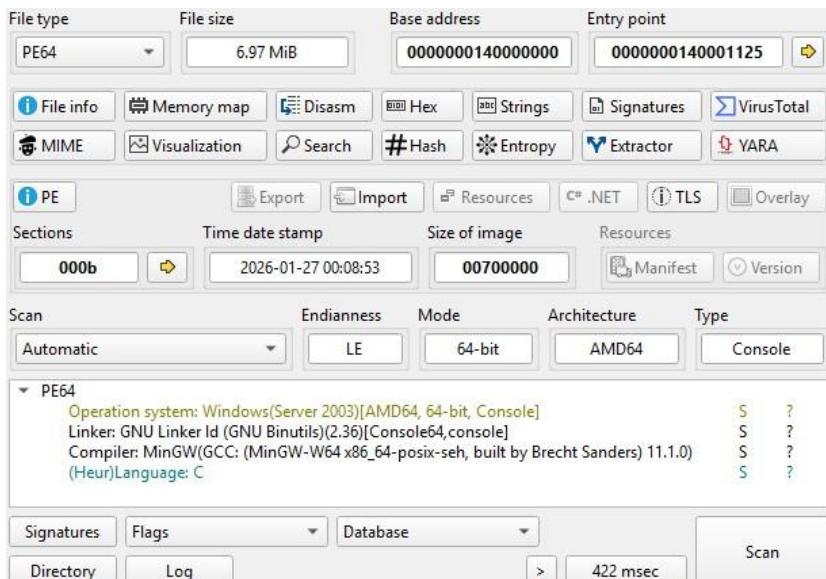
**Flag:** flag{FSM\_1s\_SO\_c00001}

**Difficulty:** Medium/Hard

**Writeup classification:** Official

## Solve

Given a single executable named `FuN51llyM4ch1n3.exe`. Running the binary executes *something* but not significant to the solve itself. We can check the binary details with DiE to see if we can find anything to note down:



Seems like a simple binary but for some reason, it's 10MB in size. Quite large for a C program.

If we attempt to execute the binary in a terminal, we get the following:

```
[REDACTED] .\FuN51llyM4ch1n3  
[REDACTED] >.\FuN51llyM4ch1n3 a  
Usage: .\FuN51llyM4ch1n3 <16 characters password>
```

Seems like it requires a 16 char password. Let's try entering an input with the length of 16 chars:

```
[REDACTED] .\FuN51llyM4ch1n3 aaaaaaaaaaaaaaaaaa  
Beep Boop.  
No.
```

Let's open it up in IDA and see what we're dealing with:

The screenshot shows the IDA Pro interface with two main panes: 'IDA View-A' (Assembly) and 'Pseudocode-A' (Pseudocode). The assembly pane displays the mainCRTStartup function, which contains numerous calls to various state\_x functions. The pseudocode pane provides a high-level overview of the logic:

```

1. void mainCRTStartup()
2. {
3.     /*(CODE) */refptr_mainCRTStartup();
4. }

```

The assembly code includes several interesting sections and comments:

- Text 0000000014000125: **S U B R O U T I N E**
- Text 0000000014000125: Attributes: bp-based frame
- Text 0000000014000125: **public: mainCRTStartup**
- Text 0000000014000125: **proc near**
- Text 0000000014000125: **mainCRTStartup**
- Text 0000000014000125: **l .DATA .XREFS : .edata:00000001140**
- Text 0000000014000125: **l .pdata:00000001404F203010**
- Text 0000000014000125: **var\_10 = dword ptr -10h**
- Text 0000000014000125: **var\_8 = dword ptr -8**
- Text 0000000014000125: **var\_4 = dword ptr -4**
- Text 0000000014000125: **arg\_0 = dword ptr 10h**
- Text 0000000014000125: **FUNCTION CHUNK AT .text:0000000240412B00 SIZE 000000220 BYTES**
- Text 0000000014000125: **l .DATA .XREFS : .edata:00000001140**
- Text 0000000014000125: **j \_unwind ( // \_C\_specific\_handler**
- Text 0000000014000125: **push rbp**
- Text 0000000014000125: **mov rbp, rsp**
- Text 0000000014000125: **sub rsp, 30h**
- Text 0000000014000125: **r mov [rbp+var\_4], RBP**
- Text 0000000014000125: **l .DATA .XREFS : .edata:00000001140**
- Text 0000000014000125: **l \_except\_l1\_start**
- Text 0000000014000125: **i \_try ( // \_except at \_l\_end**
- Text 0000000014000125: **mov rax, cs:\_refptr\_mainCRTStartup**
- Text 0000000014000125: **mov rbp, [rax+rax]**
- Text 0000000014000125: **call \_mainCRTStartup**
- Text 0000000014000125: **mov [rbp+var\_4], eax**
- Text 0000000014000125: **pop rbp**
- Text 0000000014000125: **l .DATA .XREFS : .edata:00000001140**
- Text 0000000014000125: **j // starts at 1400001134**
- Text 0000000014000125: **l .DATA .XREFS : .edata:00000001140**
- Text 0000000014000125: **i \_except\_gnu\_exception\_handler // owned by 1400001134**
- Text 0000000014000125: **l .DATA .XREFS : .edata:00000001140**
- Text 0000000014000125: **mov rax, cs:\_refptr\_g\_flag\_data**
- Text 0000000014000125: **add rbp, 30h**
- Text 0000000014000125: **pop rbp**
- Text 0000000014000125: **ret**
- Text 0000000014000125: **l .DATA .XREFS : .edata:00000001140**
- Text 0000000014000125: **j // starts at 1400001125**
- Text 0000000014000125: **l .DATA .XREFS : .edata:00000001140**
- Text 0000000014000125: **mainCRTStartup**
- Text 0000000014000125: **endp**
- Text 0000000014000125: **l .DATA .XREFS : .edata:00000001140**

In the decompilation, we see a huge function table. We also see a bunch of `state\_x` functions and in the main functions, we can see other functions such as `g\_position`, `g\_state`, `g\_transition`, and `g\_input`. Based on this, and the challenge name itself, it seems like we're dealing with some sort of Finite State Machine.

The pseudocode editor shows the mainCRTStartup function. The logic involves initializing variables (v8, g\_position, g\_transitions, g\_state), processing transitions through a state machine (using refptr\_STATES[v8]()), and printing a message if the win condition is met. The pseudocode is as follows:

```

22 v8 = (unsigned __int8)g_input ^ 0x1337LL;
23 g_position = 0;
24 g_transitions = 0;
25 g_state = v8;
26 do
27 {
28     v9 = g_position + 1;
29     if ( v8 <= 0x179FE )
30     {
31         refptr_STATES[v8]();
32         v8 = g_state;
33         v10 = g_position;
34         if ( (unsigned __int64)g_state <= 0x179FE )
35         {
36             v11 = *((__BYTE *)refptr_g_flag_data + g_state);
37             if ( v11 )
38                 flag_buf[g_position] = g_state ^ v11 ^ 0xAA;
39         }
40         v9 = v10 + 1;
41     }
42     g_position = v9;
43     --v7;
44 }
45 while ( v7 );
46 if ( g_transitions == 16 )
47 {
48     byte_1406CC0D0 = 0;
49     printf("\nBeep Boop. \nN1c3 J0b! Here: flag{%s}\n", flag_buf);
50 }
51 else
52 {
53     printf("\nBeep Boop. \nNo.\n");
54 }
55 return 0;
56 }
57 else
58 {
59     printf("Usage: .\\FuN5illyM4ch1n3 <16 characters password>\n");
60 }
61 }
62 }

```

We can see in the pseudocode that the win condition is just to provide a 16-character string where each character causes a valid transition in the state machine. Each of those 16 transitions must successfully increment `g\_transitions`.



We can also see that the starting state is calculated by XORing the first character of our input with 0x1337 and taking the modulo of the total state space (96,767).

`g_state = (0x1337 ⊕ input[0])(mod 96767)`

In each iteration of the program, it executes a function pointer from a massive array called `STATES` based on the current `g\_state`.

```
1 .data:00000014041A010 deregister_frame_fn dq 0          ; DATA XREF: __gcc_register_frame+5B1W
2 .data:00000014041A018
3 .data:00000014041A018 align 20h
4 .data:00000014041A020 public STATES
5 .data:00000014041A020 STATES dq offset state_0          ; DATA XREF: .rdata:_refptr_STATES!o
6 .data:00000014041A028 dq offset state_1
7 .data:00000014041A030 dq offset state_2
8 .data:00000014041A038 dq offset state_3
9 .data:00000014041A040 dq offset state_4
10 .data:00000014041A048 dq offset state_5
11 .data:00000014041A050 dq offset state_6
12 .data:00000014041A058 dq offset state_7
13 .data:00000014041A060 dq offset state_8
14 .data:00000014041A068 dq offset state_9
15 .data:00000014041A070 dq offset state_10
16 .data:00000014041A078 dq offset state_11
17 .data:00000014041A080 dq offset state_12
18 .data:00000014041A088 dq offset state_13
19 .data:00000014041A090 dq offset state_14
20 .data:00000014041A098 dq offset state_15
21 .data:00000014041A0A0 dq offset state_16
22 .data:00000014041A0A8 dq offset state_17
23 .data:00000014041A0B0 dq offset state_18
24 .data:00000014041A0B8 dq offset state_19
25 .data:00000014041A0C0 dq offset state_20
26 .data:00000014041A0C8 dq offset state_21
27 .data:00000014041A0D0 dq offset state_22
28 .data:00000014041A0D8 dq offset state_23
29 .data:00000014041A0E0 dq offset state_24
30 .data:00000014041A0E8 dq offset state_25
31 .data:00000014041A0F0 dq offset state_26
32 .data:00000014041A0F8 dq offset state_27
33 .data:00000014041A100 dq offset state_28
34 .data:00000014041A108 dq offset state_29
35 .data:00000014041A110 dq offset state_30
36 .data:00000014041A118 dq offset state_31
37 .data:00000014041A120 dq offset state_32
38 .data:00000014041A128 dq offset state_33
39 .data:00000014041A130 dq offset state_34
40 .data:00000014041A138 dq offset state_35
41 .data:00000014041A140 dq offset state_36
```

## State Machine Mechanics

The core of this challenge lies in the `STATES` array. Each entry is a pointer to a unique function representing a specific state.

When we analyze one of the states, we see the following:

```
1 int64 state_16()
2 {
3     char current_input; // a1
4
5     current_input = get_current_input();
6     if ( current_input == 107 )
7     {
8         increment_transitions();
9         return set_next_state(9878);
10    }
11   else if ( current_input == 74 )
12   {
13       increment_transitions();
14       return set_next_state(28202);
15   }
16   else
17   {
18       return show_error_message();
19   }
20 }
```



We see that in this state, it fetches the character at the current input position using `get\_current\_input()`, then compares that character against one or more hardcoded values. If it's a match, then it calls `increment\_transitions()` and updates `g\_state` to a new ID via `set\_next\_state()`. If it fails: It calls `show\_error\_message()`, which acts as a dead end.

To break it down in lament terms:

In `state\_16`, if the current character is 'k', we move to state 9878. But if it is 'J', we move to 28202.

Obviously, we're not going to explore the 90,000+ states manually, instead we'll utilize the Breath-First-Search to walk the graph for us. We can do this by crafting a script to achieve this.

## Locating the State Machine

The core of the binary is a massive array of function pointers named STATES located at 0x14041a020. Each pointer in this array leads to a unique function representing a specific "state" in the machine.

- Extraction: Read 96,767 quadwords from this address to map every state ID to its corresponding memory address.
- Initialization: The program determines the starting state by XORing the first character of the input with 0x1337 and taking the modulo of the total number of states.

```
import idc
import idaapi
import idautils
from collections import deque

STATES_ARRAY_ADDR = 0x14041a020
NUM_STATES = 96767

def calculate_entry_state(first_char):
    return (0x1337 ^ ord(first_char)) % NUM_STATES

def extract_addresses():
    print(f"[*] Reading {NUM_STATES} pointers from 0x{STATES_ARRAY_ADDR:x}")
    addrs
= []
    ea = STATES_ARRAY_ADDR
    for i in
range(NUM_STATES):           if i
% 10000 == 0:
        print(f"[*] {i}/{NUM_STATES}...")

        ptr =
idc.get_qword(ea)
addrs.append(ptr)           ea +=
8

    return addrs
```





## Dissecting the state transitions

Every state function follows a rigid pattern to decide which state comes next based on our input. Instead of reading these manually, we use the script to parse the assembly instructions:

- Input Comparison: The script scans for cmp instructions where the input character (stored in al or cl) is checked against a hardcoded value.
- Success Branching: It follows the jz (jump if zero) or je (jump if equal) instructions to find the logic that executes when a character match is found.
- Destination Extraction: It looks for a mov instruction (typically into registers like ecx or rdi) that sets the next state ID, allowing us to draw an "edge" between states.

```
def parse_small_state(func_ea, state_id):    if not
idc.get_func_attr(func_ea, idc.FUNCATTR_START):
    idc.create_function(func_ea)
    idc.set_name(func_ea, f"state_{state_id}", idc.SN_NOWARN)

    transitions = {}          func_start = idc.get_func_attr(func_ea,
idc.FUNCATTR_START)           func_end   = idc.get_func_attr(func_ea,
idc.FUNCATTR_END)
    if not func_end or func_end == idc.BADADDR:
        func_end = func_start + 0x200

    ea = func_start
max_ins = 100      count = 0
    current_char = None
    while ea < func_end and count < max_ins:
        mnem = idc.print_insn_mnem(ea)
        if mnem ==
"cmp":
            op1 = idc.print_operand(ea, 0)
            if op1 in ["al",
"cl"]:
                char_val = idc.get_operand_value(ea, 1)
if 0x20 <= char_val <= 0x7E:
                current_char = chr(char_val)
```

```

        elif mnem in ["jz", "je"] and current_char:
target = idc.get_operand_value(ea, 0)

        scan_ea = target
for _ in range(10):
        if
scan_ea >= func_end:
            break
        scan_mnem =
idc.print_insn_mnem(scan_ea)
        if
scan_mnem == "mov":
            scan_op1 = idc.print_operand(scan_ea, 0)
            if scan_op1 in ["ecx",
"rcx", "edi", "rdi"]:
                next_state = idc.get_operand_value(scan_ea, 1)
                if 0 <=
next_state < NUM_STATES:
                    transitions[current_char] =
next_state
                    current_char = None
                break
            scan_ea =
idc.next_head(scan_ea)
            elif mnem
== "call":
                target_name = idc.get_func_name(idc.get_operand_value(ea, 0))
if target_name == "get_current_input":
                current_char = None

            ea      =      idc.next_head(ea)
count += 1

return transitions

```

## Breath-First-Search Solve

Because the state space is too large for a simple brute-force approach, we use a BreadthFirst Search (BFS) to "walk" the graph efficiently.

Since the first character of the password determines the starting state and acts as the first transition, we iterate through all printable characters (A-Z, 0-9):

- Calculate the `entry\_state` for a character.
- Parse that `entry\_state` function to see if it accepts that same character.
- If it does, we have a valid starting point for our search.

```
def build_graph(addr):
    print("[*] Building graph...")
    graph = {}
    visited = set()

    import string
    queue = deque()
        for first_char in string.ascii_letters +
string.digits:
            entry      = calculate_entry_state(first_char)
    if entry not in visited and entry < len(addr):
        queue.append(entry)
        print(f"[*] Starting BFS from {len(queue)} potential entry
states...")
        while
queue:
            sid = queue.popleft()
                if sid in visited or sid >=
len(addr):
                    continue
            visited.add(sid)
                if len(visited) %
100 == 0:
                    print(f"[*] States: {len(visited)}", end='\r')

            trans = parse_small_state(addr[sid], sid)
    graph[sid] = trans
```

```
if Len(visited) <= 10:
    print(f"\n[DEBUG] State {sid} at 0x{addrs[sid]:x}:
{Len(trans)} transitions")
    for c, n in trans.items():
print(f"  '{c}' -> {n}")
        for nxt in
trans.values():
            if nxt
not in visited:
queue.append(nxt)

print(f"\n[+] Graph: {Len(graph)} states, {sum(Len(t) for t in
graph.values())} edges")
return graph
```

# Graph Traversal and Pruning

From a valid starting point, the BFS explores every possible valid transition layer-by-layer.

- Path Length: We only care about paths that reach exactly 16 characters.
- Visited States: To prevent the script from getting stuck in infinite loops, we track (state\_id, path\_length). If we reach the same state at the same character position again, we discard that branch.

```
def solve(graph):
    print("[*] Solving for 16-character password...")

    import string
    solutions = []
        for fc in string.ascii_letters +
string.digits:
            entry = calculate_entry_state(fc)
            if entry not
in graph:
            continue
            if fc not in
graph[entry]:
                continue
                print(f"[*] Trying '{fc}' (entry={entry})...", end='\r')

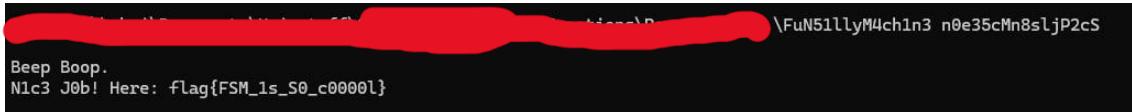
            first_next = graph[entry][fc]
            q = deque([(first_next, fc)])
            vis = set()
                while q and
len(solutions) < 10:
                    s, p = q.popleft()
                    if (s,
len(p)) in vis:
                        continue
                        vis.add((s,
len(p)))
                        if
len(p) == 16:
                            print(f"\n[+] SOLUTION: {p}")
                            solutions.append(p)
                            break
```

```
        if s in graph:  
for c, n in graph[s].items():  
if len(p) < 16:  
            q.append((n, p + c))  
return solutions
```

Running this script in IDA > File > Script file, we're able to find the exact password that the program uses:

```
[DEBUG] State 4957 at 0x14037d4b0: 0 transitions  
[*] States: 100  
[+] Graph: 146 states, 84 edges  
[*] Solving for 16-character password...  
[*] Trying 'n' (entry=4953)...  
[+] SOLUTION: n0e35cMn8s1jP2cS  
  
=====  
[+] Found 1 solution(s):  
n0e35cMn8s1jP2cS
```

Inputting this password into the program, we get the flag:



Flag = flag{FSM\_1s\_S0\_c00001}